

# demo\_moon

December 4, 2025

```
[1]: from sklearn.datasets import make_moons
import torch
from torch import nn
from torch.utils.data import DataLoader, TensorDataset
from tqdm import tqdm
import os , sys
sys.path.append(os.path.abspath(os.path.join(os.getcwd(), '..')))
from core.optimization import NSGACConfig, run_nsga
from core.cf_problem import make_cf_problem
from core.data import DataGenerator, DatasetsDG
from core.models import SimpleNN, EnsembleModel
from matplotlib import pyplot as plt
from utils import plot_proba, plot_uncertainty_heatmap, plot_cf_3d
```


```
[2]: import importlib
import core.optimization
import core.cf_problem

importlib.reload(core.optimization)
importlib.reload(core.cf_problem)
```

```
[2]: <module 'core.cf_problem' from 'e:\\Thieses mit
Thies\\uamocf\\core\\cf_problem.py'>
```

```
[3]: ds=make_moons(n_samples=10000, noise=0.5, random_state=42)

tds = TensorDataset(torch.tensor(ds[0], dtype=torch.float32), torch.
    tensor(ds[1], dtype=torch.long))
dg=DatasetsDG(tds, num_classes=2)
dg.plot()
```



```
[4]: models = [SimpleNN(input_dim=2, hidden_dim=14, output_dim=2, depth=3) for _ in range(5)]
ensemble_model = EnsembleModel(models=models)

# train each model in the ensemble
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
ensemble_model.to(device)
epochs = 150
val_data = dg.sample(500, seed=123)
for model in ensemble_model.models:
    seed = 42 + hash(model) % 1000 # Simple way to get different seeds
    samples = dg.sample(400, seed=seed)
    _tds= TensorDataset(samples[0].to(device), samples[1].to(device))
    DL=DataLoader(_tds, batch_size=256, shuffle=True)
    model.train()
```

```

bar = tqdm(range(epochs), desc="Training Model", colour="blue", unit="epoch")
running_loss = 0.0
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()
for epoch in bar:
    for inputs, labels in DL:
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
    bar.set_postfix({"loss": running_loss / (epoch + 1)})

model.eval()
with torch.no_grad():
    val_inputs, val_labels = val_data[0].to(device), val_data[1].to(device)
    val_outputs = model(val_inputs)
    _, val_preds = torch.max(val_outputs, 1)
    _, val_labels_idx = torch.max(val_labels, 1)
    val_accuracy = (val_preds == val_labels_idx).float().mean().item()
    print(f"Validation Accuracy: {val_accuracy * 100:.2f}%")
print(f"Finished Training Model")

```

Training Model: 100% | 150/150 [00:02<00:00, 50.07epoch/s, loss=0.94]

Validation Accuracy: 80.20%

Finished Training Model

Training Model: 100% | 150/150 [00:01<00:00, 120.15epoch/s, loss=0.933]

Validation Accuracy: 80.80%

Finished Training Model

Training Model: 100% | 150/150 [00:01<00:00, 113.39epoch/s, loss=0.92]

Validation Accuracy: 79.80%

Finished Training Model

Training Model: 100% | 150/150 [00:01<00:00, 115.33epoch/s, loss=0.955]

Validation Accuracy: 80.40%

Finished Training Model

Training Model: 100% | 150/150 [00:01<00:00, 117.83epoch/s, loss=0.975]

Validation Accuracy: 79.40%  
Finished Training Model

[5]: # plot decision boundary


```
fig, axes = plt.subplots(1, 3, figsize=(18, 5))

plot_proba(ensemble_model.models[0], ds[0], levels=5, ax=axes[0])
axes[0].set_title("Decision Boundary")

plot_uncertainty_heatmap(ensemble_model, ds[0], uncertainty_type="epistemic", ↴device=device, levels=5, ax=axes[1])
axes[1].set_title("Epistemic Uncertainty")

plot_uncertainty_heatmap(ensemble_model, ds[0], uncertainty_type="aleatoric", ↴device=device, levels=5, ax=axes[2])
axes[2].set_title("Aleatoric Uncertainty")

plt.tight_layout()
plt.show()
```



[6]:

```
x_factual = torch.tensor([[2, -1]], dtype=torch.float32)
y_target = torch.tensor([0], dtype=torch.long)
X_obs = dg.sample(1000, seed=456)[0]
if isinstance(X_obs, torch.Tensor):
    X_obs = X_obs.clone().detach().float()
else:
    X_obs = torch.tensor(X_obs, dtype=torch.float32)

NSGA2_config = NSGAConfig(
    pop_size=150,
    min_gen=250,
    max_gen=5950,
```

```

        use_conditional_mutator=True,
        conditional_mutator_prob=0.7,
        use_reset_operator=True,
        reset_prob=0.05,
    )

# ===== VERSION 1: With Uncertainty (EU/AU) =====
print("==" * 50)
print("Running VERSION 1: With Epistemic & Aleatoric Uncertainty")
print("Objectives: [Validity, EU, Sparsity, -AU]")
print("==" * 50)
problem_uncertainty = make_cf_problem(
    ensemble_model.models[0], x_factual, y_target, X_obs,
    device=device,
    ensemble=ensemble_model # Enables EU/AU objectives
)
results_uncertainty = run_nsga(problem_uncertainty, NSGA2_config, X_obs=X_obs.
    ↪numpy(), x_factual=x_factual.numpy())

# ===== VERSION 2: Simple (Similarity/Plausibility) =====
print("\n" + "==" * 50)
print("Running VERSION 2: Simple (Similarity & Plausibility)")
print("Objectives: [Validity, Similarity, Sparsity, Plausibility]")
print("==" * 50)
problem_simple = make_cf_problem(
    ensemble_model.models[0], x_factual, y_target, X_obs,
    device=device,
    ensemble=None # No ensemble = uses Similarity/Plausibility objectives
)
results_simple = run_nsga(problem_simple, NSGA2_config, X_obs=X_obs.numpy(),
    ↪x_factual=x_factual.numpy())

print("\n Both optimizations completed!")

```

```

=====
Running VERSION 1: With Epistemic & Aleatoric Uncertainty
Objectives: [Validity, EU, Sparsity, -AU]
=====
Gen   1 | Valid CFs (pop):  66 | Pareto front:  48 | Best P(target): 0.994 |
Mean Sparsity: 1.979
Gen   10 | Valid CFs (pop): 124 | Pareto front: 150 | Best P(target): 0.993 |
Mean Sparsity: 1.773
Gen   20 | Valid CFs (pop): 128 | Pareto front: 150 | Best P(target): 0.994 |
Mean Sparsity: 1.787
Gen   30 | Valid CFs (pop): 132 | Pareto front: 150 | Best P(target): 0.995 |
Mean Sparsity: 1.800
Gen   40 | Valid CFs (pop): 124 | Pareto front: 150 | Best P(target): 0.994 |

```

Mean Sparsity: 1.793  
 Gen 50 | Valid CFs (pop): 129 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.807  
 Gen 60 | Valid CFs (pop): 132 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.787  
 Gen 70 | Valid CFs (pop): 130 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.807  
 Gen 80 | Valid CFs (pop): 127 | Pareto front: 150 | Best P(target): 0.995 |  
 Mean Sparsity: 1.767  
 Gen 90 | Valid CFs (pop): 128 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.773  
 Gen 100 | Valid CFs (pop): 126 | Pareto front: 150 | Best P(target): 0.995 |  
 Mean Sparsity: 1.767  
 Gen 110 | Valid CFs (pop): 130 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.787  
 Gen 120 | Valid CFs (pop): 127 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.767  
 Gen 130 | Valid CFs (pop): 127 | Pareto front: 150 | Best P(target): 0.995 |  
 Mean Sparsity: 1.793  
 Gen 140 | Valid CFs (pop): 130 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.807  
 Gen 150 | Valid CFs (pop): 131 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.780  
 Gen 160 | Valid CFs (pop): 130 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.787  
 Gen 170 | Valid CFs (pop): 128 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.780  
 Gen 180 | Valid CFs (pop): 125 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.760  
 Gen 190 | Valid CFs (pop): 132 | Pareto front: 150 | Best P(target): 0.994 |  
 Mean Sparsity: 1.813  
 Gen 200 | Valid CFs (pop): 124 | Pareto front: 150 | Best P(target): 0.995 |  
 Mean Sparsity: 1.720  
 Gen 210 | Valid CFs (pop): 132 | Pareto front: 150 | Best P(target): 0.995 |  
 Mean Sparsity: 1.787  
 Gen 220 | Valid CFs (pop): 126 | Pareto front: 150 | Best P(target): 0.995 |  
 Mean Sparsity: 1.753  
 Gen 230 | Valid CFs (pop): 130 | Pareto front: 150 | Best P(target): 0.995 |  
 Mean Sparsity: 1.833  
 Gen 240 | Valid CFs (pop): 130 | Pareto front: 150 | Best P(target): 0.995 |  
 Mean Sparsity: 1.813  
 Gen 250 | Valid CFs (pop): 129 | Pareto front: 150 | Best P(target): 0.995 |  
 Mean Sparsity: 1.813

---

Running VERSION 2: Simple (Similarity & Plausibility)  
 Objectives: [Validity, Similarity, Sparsity, Plausibility]

---

Gen 1 | Valid CFs (pop): 66 | Pareto front: 62 | Best P(target): 0.994 |  
Mean Sparsity: 1.984  
Gen 10 | Valid CFs (pop): 107 | Pareto front: 150 | Best P(target): 0.994 |  
Mean Sparsity: 1.853  
Gen 20 | Valid CFs (pop): 99 | Pareto front: 150 | Best P(target): 0.993 |  
Mean Sparsity: 1.853  
Gen 30 | Valid CFs (pop): 110 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.867  
Gen 40 | Valid CFs (pop): 105 | Pareto front: 150 | Best P(target): 0.994 |  
Mean Sparsity: 1.840  
Gen 50 | Valid CFs (pop): 105 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.873  
Gen 60 | Valid CFs (pop): 104 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.833  
Gen 70 | Valid CFs (pop): 102 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.833  
Gen 80 | Valid CFs (pop): 111 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.873  
Gen 90 | Valid CFs (pop): 114 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.847  
Gen 100 | Valid CFs (pop): 104 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.867  
Gen 110 | Valid CFs (pop): 106 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.860  
Gen 120 | Valid CFs (pop): 98 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.807  
Gen 130 | Valid CFs (pop): 104 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.847  
Gen 140 | Valid CFs (pop): 104 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.813  
Gen 150 | Valid CFs (pop): 106 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.847  
Gen 160 | Valid CFs (pop): 103 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.833  
Gen 170 | Valid CFs (pop): 103 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.853  
Gen 180 | Valid CFs (pop): 105 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.847  
Gen 190 | Valid CFs (pop): 114 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.887  
Gen 200 | Valid CFs (pop): 107 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.787  
Gen 210 | Valid CFs (pop): 112 | Pareto front: 150 | Best P(target): 0.994 |  
Mean Sparsity: 1.833  
Gen 220 | Valid CFs (pop): 101 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.847  
Gen 230 | Valid CFs (pop): 106 | Pareto front: 150 | Best P(target): 0.995 |  
Mean Sparsity: 1.840

```

Gen 240 | Valid CFs (pop): 108 | Pareto front: 150 | Best P(target): 0.994 |
Mean Sparsity: 1.867
Gen 250 | Valid CFs (pop): 103 | Pareto front: 150 | Best P(target): 0.994 |
Mean Sparsity: 1.840

```

Both optimizations completed!

```
[7]: # Compare both versions in 2D feature space
fig, axes = plt.subplots(2, 3, figsize=(18, 10))

cf_uncertainty = torch.tensor(results_uncertainty.X, dtype=torch.float32).cpu().numpy()
cf_simple = torch.tensor(results_simple.X, dtype=torch.float32).cpu().numpy()

# ===== ROW 1: VERSION 1 - With Uncertainty (EU/AU) =====
# 1. Probability background
plot_proba(ensemble_model.models[0], ds[0], levels=5, ax=axes[0, 0])
axes[0, 0].scatter(cf_uncertainty[:, 0], cf_uncertainty[:, 1], color='red', marker='x', s=100, label='CFs')
axes[0, 0].scatter(x_factual[0, 0].item(), x_factual[0, 1].item(), color='green', marker='o', s=100, label='Factual')
axes[0, 0].set_title("V1 (EU/AU): Probability")
axes[0, 0].legend()

# 2. Epistemic Uncertainty background
plot_uncertainty_heatmap(ensemble_model, ds[0], device=device, uncertainty_type="epistemic", levels=5, ax=axes[0, 1])
axes[0, 1].scatter(cf_uncertainty[:, 0], cf_uncertainty[:, 1], color='red', marker='x', s=100, label='CFs')
axes[0, 1].scatter(x_factual[0, 0].item(), x_factual[0, 1].item(), color='green', marker='o', s=100, label='Factual')
axes[0, 1].set_title("V1 (EU/AU): Epistemic Uncertainty")
axes[0, 1].legend()

# 3. Aleatoric Uncertainty background
plot_uncertainty_heatmap(ensemble_model, ds[0], device=device, uncertainty_type="aleatoric", levels=5, ax=axes[0, 2])
axes[0, 2].scatter(cf_uncertainty[:, 0], cf_uncertainty[:, 1], color='red', marker='x', s=100, label='CFs')
axes[0, 2].scatter(x_factual[0, 0].item(), x_factual[0, 1].item(), color='green', marker='o', s=100, label='Factual')
axes[0, 2].set_title("V1 (EU/AU): Aleatoric Uncertainty")
axes[0, 2].legend()

# ===== ROW 2: VERSION 2 - Simple (Similarity/Plausibility) =====
# 1. Probability background
plot_proba(ensemble_model.models[0], ds[0], levels=5, ax=axes[1, 0])

```

```


axes[1, 0].scatter(cf_simple[:, 0], cf_simple[:, 1], color='blue', marker='x', s=100, label='CFs')
axes[1, 0].scatter(x_factual[0, 0].item(), x_factual[0, 1].item(), color='green', marker='o', s=100, label='Factual')
axes[1, 0].set_title("V2 (Simple): Probability")
axes[1, 0].legend()

# 2. Epistemic Uncertainty background (for comparison)
plot_uncertainty_heatmap(ensemble_model, ds[0], device=device,
    uncertainty_type="epistemic", levels=5, ax=axes[1, 1])
axes[1, 1].scatter(cf_simple[:, 0], cf_simple[:, 1], color='blue', marker='x', s=100, label='CFs')
axes[1, 1].scatter(x_factual[0, 0].item(), x_factual[0, 1].item(), color='green', marker='o', s=100, label='Factual')
axes[1, 1].set_title("V2 (Simple): on EU background")
axes[1, 1].legend()

# 3. Aleatoric Uncertainty background (for comparison)
plot_uncertainty_heatmap(ensemble_model, ds[0], device=device,
    uncertainty_type="aleatoric", levels=5, ax=axes[1, 2])
axes[1, 2].scatter(cf_simple[:, 0], cf_simple[:, 1], color='blue', marker='x', s=100, label='CFs')
axes[1, 2].scatter(x_factual[0, 0].item(), x_factual[0, 1].item(), color='green', marker='o', s=100, label='Factual')
axes[1, 2].set_title("V2 (Simple): on AU background")
axes[1, 2].legend()

plt.suptitle("Comparison: V1 (with EU/AU objectives) vs V2 (Similarity/  
Plausibility)", fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()

```



```
[8]: import importlib
import utils
importlib.reload(utils)
from utils import plot_cf_3d

# Convert tensors to numpy arrays for the plot
x_factual_np = x_factual.cpu().numpy() if isinstance(x_factual, torch.Tensor) else x_factual
context_data = dg.sample(10000, seed=789)[0]
context_np = context_data.cpu().numpy() if isinstance(context_data, torch.Tensor) else context_data

# ===== 3D Objective Space: VERSION 1 (EU/AU) =====
print("==" * 50)
print("3D Objective Space: VERSION 1 (with EU/AU)")
print("Axes: EU (x), Validity (y), AU (z - reversed)")
print("==" * 50)
plot_cf_3d(results=results_uncertainty, x_factual=x_factual_np, context=context_np)

=====
```

```
3D Objective Space: VERSION 1 (with EU/AU)
Axes: EU (x), Validity (y), AU (z - reversed)
=====
```

```
[9]: # ===== 3D Objective Space: VERSION 2 (Simple) =====
print("==" * 50)
```

```

print("3D Objective Space: VERSION 2 (Similarity/Plausibility)")
print("Axes: Similarity (x), Validity (y), Plausibility (z)")
print("=" * 50)
plot_cf_3d(results=results_simple, x_factual=x_factual_np, context=context_np)

```

```

=====
3D Objective Space: VERSION 2 (Similarity/Plausibility)
Axes: Similarity (x), Validity (y), Plausibility (z)
=====
```

[10]: # ===== Summary Comparison =====

```

print("=" * 60)
print("SUMMARY: Objective Values Comparison")
print("=" * 60)

print("\n--- VERSION 1: With Uncertainty (EU/AU) ---")
print("Objectives: [Validity, Epistemic U., Sparsity, -Aleatoric U.]")
print(f"Number of Pareto solutions: {len(results_uncertainty.X)}")
print(f"Validity range: [{results_uncertainty.F[:,0].min():.4f},"
    ↪{results_uncertainty.F[:,0].max():.4f}]")
print(f"Epistemic U. range: [{results_uncertainty.F[:,1].min():.4f},"
    ↪{results_uncertainty.F[:,1].max():.4f}]")
print(f"Sparsity range: [{results_uncertainty.F[:,2].min():.4f},"
    ↪{results_uncertainty.F[:,2].max():.4f}]")
print(f"-AU range: [{results_uncertainty.F[:,3].min():.4f},"
    ↪{results_uncertainty.F[:,3].max():.4f}]")
print(f"→ AU (positive): [{-results_uncertainty.F[:,3].max():.4f},"
    ↪{-results_uncertainty.F[:,3].min():.4f}]")

print("\n--- VERSION 2: Simple (Similarity/Plausibility) ---")
print("Objectives: [Validity, Similarity, Sparsity, Plausibility]")
print(f"Number of Pareto solutions: {len(results_simple.X)}")
print(f"Validity range: [{results_simple.F[:,0].min():.4f}, {results_simple.F[:",
    ↪,0].max():.4f}]")
print(f"Similarity range: [{results_simple.F[:,1].min():.4f}, {results_simple.F[:",
    ↪,1].max():.4f}]")
print(f"Sparsity range: [{results_simple.F[:,2].min():.4f}, {results_simple.F[:",
    ↪,2].max():.4f}]")
print(f"Plausibility range: [{results_simple.F[:,3].min():.4f}, {results_simple.F[:",
    ↪,3].max():.4f}]")

```

```

=====
SUMMARY: Objective Values Comparison
=====
```

```

--- VERSION 1: With Uncertainty (EU/AU) ---
Objectives: [Validity, Epistemic U., Sparsity, -Aleatoric U.]
Number of Pareto solutions: 150
```

Validity range: [0.0048, 0.9890]  
Epistemic U. range: [0.0001, 0.0426]  
Sparsity range: [0.0000, 2.0000]  
-AU range: [-0.6926, -0.0487]  
→ AU (positive): [0.0487, 0.6926]

--- VERSION 2: Simple (Similarity/Plausibility) ---  
Objectives: [Validity, Similarity, Sparsity, Plausibility]  
Number of Pareto solutions: 150  
Validity range: [0.0058, 0.9890]  
Similarity range: [0.0000, 0.7714]  
Sparsity range: [0.0000, 2.0000]  
Plausibility range: [0.0058, 0.1270]