

nd_mmo

November 23, 2025

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import torch
from torch import nn, optim
from torch.utils.data import TensorDataset, DataLoader, random_split

from pymoo.algorithms.moo.nsga2 import NSGA2
from pymoo.optimize import minimize
from pymoo.termination import get_termination
from pymoo.problems.functional import FunctionalProblem

from tqdm import tqdm

from sklearn.datasets import make_moons
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.svm import SVC
from sklearn.inspection import DecisionBoundaryDisplay
```

```
[70]: import numpy as np
import matplotlib.pyplot as plt

# -----
# 1) Boundary-focused probabilistic boundary
# -----

def decision_prob_boundary_focus(
    x, y,
    r_inner=np.sqrt(2), r_outer=3.0,
    sigma=0.85,           # width of uncertain band around boundary
    max_uncertainty=1.0   # 1.0 -> p=0.5 exactly on boundary
):
    """
    Non-deterministic P(y=1/x) with uncertainty emphasized near boundaries.
    Hard rule:
```

```

y=1 if r<=r_inner or r>=r_outer    (inside inner disk OR outside outer_
disk)
y=0 otherwise (annulus)
Then we add boundary uncertainty:
closer to boundary -> p moves toward 0.5
far away -> p -> hard label (0 or 1)
"""

r = np.sqrt(x**2 + y**2)

# hard/deterministic label
base = 1.0 if (r <= r_inner or r >= r_outer) else 0.0

# distance to nearest boundary
dist_to_boundary = min(abs(r - r_inner), abs(r - r_outer))

# uncertainty weight peaked at boundary (Gaussian bump)
alpha = max_uncertainty * np.exp(-(dist_to_boundary / sigma)**2)

# blend base toward 0.5 near boundary
p = base * (1 - alpha) + 0.5 * alpha
return float(np.clip(p, 0.0, 1.0))

def generate_data(n=300, prob_fn=decision_prob_boundary_focus, xlim=(-4, 4),
rng=None):
"""
Sample X uniformly, then sample Y ~ Bernoulli(p).
Returns X, Y, p1 (probability for class 1).
"""

rng = np.random.default_rng() if rng is None else rng
X = rng.uniform(xlim[0], xlim[1], size=(n, 2))
p1 = np.array([prob_fn(x, y) for x, y in X])
Y = rng.binomial(1, p1).astype(np.int64)
return X, Y, p1

# -----
# 2) Gradient visualization of P(y=1/x)
# -----
def plot_probabilistic_boundary(prob_fn, xlim=(-4, 4), grid_n=450):
    xs = np.linspace(xlim[0], xlim[1], grid_n)
    ys = np.linspace(xlim[0], xlim[1], grid_n)
    xx, yy = np.meshgrid(xs, ys)

    # evaluate probability field
    pp = np.vectorize(prob_fn)(xx, yy)

```

```

plt.figure(figsize=(6.5, 5.5))
#   color map showing probability
cf = plt.contourf(xx, yy, pp, levels=80, cmap="viridis")
plt.colorbar(cf, label="P(y=1 | x)")

# emphasize the "closing decision boundary area" by drawing p=0.5 contour
plt.contour(xx, yy, pp, levels=[0.5], colors="white", linewidths=2)

plt.title("Non-deterministic decision boundary (uncertainty near boundary)")
plt.xlabel("x"); plt.ylabel("y")
plt.show()


# ---- run / demo ----
plot_probabilistic_boundary(decision_prob_boundary_focus)


X, Y, p1 = generate_data(300)

plt.figure(figsize=(6.5, 5.5))
plt.scatter(X[:, 0], X[:, 1], c=p1, s=25, cmap="viridis")
plt.colorbar(label="P(y=1 | x) on samples")
plt.title("Sampled data colored by probability ( )")
plt.xlabel("x"); plt.ylabel("y")
plt.show()

```

Non-deterministic decision boundary (uncertainty near boundary)





```
[71]: # -----
# 2) Train classifier
# -----
device = "cuda" if torch.cuda.is_available() else "cpu"

dataset = TensorDataset(torch.tensor(X, dtype=torch.float32),
                      torch.tensor(Y, dtype=torch.long))

train_size = int(0.8 * len(dataset))
val_size = len(dataset) - train_size
train_ds, val_ds = random_split(dataset, [train_size, val_size])

train_loader = DataLoader(train_ds, batch_size=32, shuffle=True)
val_loader = DataLoader(val_ds, batch_size=256, shuffle=False)

class SimpleNN(nn.Module):
```

```

def __init__(self):
    super().__init__()
    self.net = nn.Sequential(
        nn.Linear(2, 16), nn.ReLU(),
        nn.Linear(16, 16), nn.ReLU(),
        nn.Linear(16, 16), nn.ReLU(),
        nn.Linear(16, 16), nn.ReLU(),
        nn.Linear(16, 2)
    )

    def forward(self, x):
        return self.net(x)

model = SimpleNN().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)

```

```

[72]: # training loop
num_epochs = 100

bar = tqdm(range(num_epochs), desc="Training", ncols=80, unit="epoch", colour="blue")
for epoch in bar:
    model.train()
    running_loss = 0.0
    for inputs, labels in train_loader:
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * inputs.size(0)

    epoch_loss = running_loss / len(train_loader.dataset)
    bar.set_postfix({'Train Loss': epoch_loss})

# Validation
model.eval()
val_loss = 0.0
correct = 0
with torch.no_grad():
    for inputs, labels in val_loader:
        inputs, labels = inputs.to(device), labels.to(device)

```

```

outputs = model(inputs)
loss = criterion(outputs, labels)
val_loss += loss.item() * inputs.size(0)
_, preds = torch.max(outputs, 1)
correct += (preds == labels).sum().item()

val_loss /= len(val_loader.dataset)
val_accuracy = correct / len(val_loader.dataset)

bar.set_postfix({'Train Loss': epoch_loss, 'Val Loss': val_loss, 'Val Acc': val_accuracy})

```

Training: 100%| 100/100 [00:02<00:00, 36.89epoch/s, Train Loss=0.481]

```

[73]: import numpy as np
import matplotlib.pyplot as plt
from matplotlib import cm
from sklearn.inspection import DecisionBoundaryDisplay

def plot_torch_proba_boundary(
    torch_model,
    X,
    y=None,
    class_of_interest="auto",      # "auto" / None/"max" / int / "all"
    grid_resolution=300,
    padding=0.8,
    levels=100,
    cmap_single="Blues",
    scatter_kwarg=None,
    ax=None,
    device=None,
    already_prob=False,
    multiclass_cmap="tab10",       # <--- NEW: background+points share this
):
    X = np.asarray(X)
    assert X.shape[1] == 2, "X must be shape (N,2)"

    est = TorchProbaEstimator(
        torch_model,
        device=device,
        already_prob=already_prob
    )
    est.fit(X, y)
    n_classes = len(est.classes_)

    # --- Build the exact class color palette used for background regions ---

```

```

base_cmap = plt.cm.get_cmap(multiclass_cmap, n_classes)
multiclass_colors = base_cmap(np.arange(n_classes)) # (n_classes,4 RGBA)

if scatter_kw_args is None:
    scatter_kw_args = dict(
        s=28, marker="o",
        linewidths=0.9, edgecolor="k",
        alpha=0.9
    )

if class_of_interest == "auto":
    class_of_interest = 1 if n_classes == 2 else None

y_arr = None if y is None else np.asarray(y)

# ---- multi-class: each class prob subplot ----
if class_of_interest == "all" and n_classes > 2:
    if ax is None:
        fig, axes = plt.subplots(
            1, n_classes, figsize=(4.2*n_classes, 3.6),
            constrained_layout=True
        )
    else:
        axes = ax
        fig = axes[0].figure

    for k in range(n_classes):
        disp = DecisionBoundaryDisplay.from_estimator(
            est,
            X,
            response_method="predict_proba",
            class_of_interest=k,
            plot_method="contourf",
            levels=levels,
            vmin=0, vmax=1,
            cmap=cmap_single,
            ax=axes[k],
            grid_resolution=grid_resolution,
            eps=padding,
        )
        axes[k].set_title(f"Class {k} probability")
        axes[k].set(xticks=(), yticks=())

    # points colored by TRUE class using the SAME class palette
    if y_arr is not None:
        for kk in range(n_classes):
            mask = y_arr == est.classes_[kk]

```

```

        axes[k].scatter(
            X[mask, 0], X[mask, 1],
            color=multiclass_colors[kk],
            **scatter_kwargs
        )

    fig.colorbar(
        disp.surface_, ax=axes, orientation="horizontal",
        fraction=0.05, pad=0.08, label="Probability"
    )
    return fig, axes

# ---- single plot ----
if ax is None:
    fig, ax = plt.subplots(1, 1, figsize=(5.2, 4.6))
else:
    fig = ax.figure

# ---- max-class regions (multi-class background) ----
if class_of_interest is None or class_of_interest == "max":
    disp = DecisionBoundaryDisplay.from_estimator(
        est,
        X,
        response_method="predict_proba",
        class_of_interest=None,
        plot_method="contourf",
        levels=levels,
        vmin=0, vmax=1,
        ax=ax,
        grid_resolution=grid_resolution,
        eps=padding,
        multiclass_colors=multiclass_colors, # <-- KEY LINE
    )
    ax.set_title("Max-probability class regions")
    ax.set(xticks=(), yticks=())

# points colored by TRUE class using EXACT SAME background colors
if y_arr is not None:
    for k in range(n_classes):
        mask = y_arr == est.classes_[k]
        ax.scatter(
            X[mask, 0], X[mask, 1],
            color=multiclass_colors[k],
            **scatter_kwargs
        )

# sklearn-like per-class mini colorbars

```

```

max_cmmaps = [s.cmap for s in disp.surface_]
for k in range(n_classes):
    cax = fig.add_axes([0.78, 0.12 + k*0.06, 0.18, 0.04])
    fig.colorbar(
        cm.ScalarMappable(norm=None, cmap=max_cmmaps[k]),
        cax=cax, orientation="horizontal"
    )
    cax.set_title(f"P(class {k})", fontsize=9)

return fig, ax

# ---- probability surface for a specific class (blue heatmap) ----
k = int(class_of_interest)
disp = DecisionBoundaryDisplay.from_estimator(
    est,
    X,
    response_method="predict_proba",
    class_of_interest=k,
    plot_method="contourf",
    levels=levels,
    vmin=0, vmax=1,
    cmap=cmap_single,
    ax=ax,
    grid_resolution=grid_resolution,
    eps=padding,
)
ax.set_title(f"Class {k} probability surface")
ax.set(xticks=(), yticks=())

# still color points by TRUE class using the SAME palette as max-mode
if y_arr is not None:
    for kk in range(n_classes):
        mask = y_arr == est.classes_[kk]
        ax.scatter(
            X[mask, 0], X[mask, 1],
            color=multiclass_colors[kk],
            **scatter_kwargs
        )

fig.colorbar(disp.surface_, ax=ax, label="Probability")
return fig, ax

```


[74]: # visual

```
plot_torch_proba_boundary(model, X, levels=500)
```

```
C:\Users\yuanl\AppData\Local\Temp\ipykernel_26632\3152812757.py:33:
MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib
```

```
3.7 and will be removed in 3.11. Use ``matplotlib.colormaps[name]`` or  
``matplotlib.colormaps.get_cmap()`` or ``pyplot.get_cmap()`` instead.  
base_cmap = plt.cm.get_cmap(multiclass_cmap, n_classes)
```

```
[74]: (<Figure size 520x460 with 2 Axes>,  
<Axes: title={'center': 'Class 1 probability surface'}>)
```



```
[114]: # -----  
# 3) Multi-objective CF problem  
# -----  
def make_cf_problem(model, x_star, Y_prime, X_obs, weights):  
    model.eval()  
  
    x_star_np = x_star.detach().cpu().numpy()  
    X_obs_np = X_obs.detach().cpu().numpy()  
    w_np = weights.detach().cpu().numpy()  
    w_np = w_np / (w_np.sum() + 1e-12)  
  
    p = x_star_np.shape[0]  
  
    xl = X_obs_np.min(axis=0)  
    xu = X_obs_np.max(axis=0)  
    feature_range = xu - xl
```

```

feature_range[feature_range == 0] = 1.0

target_labels = Y_prime.view(-1).long().tolist()

def delta_G_vec(x, y):
    return np.minimum(np.abs(x - y) / feature_range, 1.0)

# 1) Validity: distance to target in probability space
def o1_validity(x):
    x_t = torch.from_numpy(x.astype(np.float32)).unsqueeze(0).to(device)
    with torch.no_grad():
        probs = torch.softmax(model(x_t), dim=1).squeeze(0).cpu().numpy()
    p_target = probs[target_labels].max()
    return float(1.0 - p_target)

# 2) Similarity
def o2_similarity(x):
    return float(delta_G_vec(x, x_star_np).mean())

# 3) Sparsity (L0)
EPS = 0.005
def o3_sparsity(x):
    diff=0
    for i in range(len(x_star_np)):
        if abs(x[i] - x_star_np[i]) > EPS:
            diff += 1
    return float(diff)

# 4) Plausibility
def o4_plausibility(x):
    diffs = (X_obs_np - x) / feature_range
    per_sample = np.sqrt((diffs**2).mean(axis=1))
    return float((per_sample * w_np).sum())

return FunctionalProblem(
    n_var=p,
    objs=[o1_validity, o2_similarity, o3_sparsity, o4_plausibility],
    xl=xl, xu=xu,
    elementwise=True
)

# pick a REAL observed instance as x*
x_star = dataset.tensors[0][0].to(device)
with torch.no_grad():

```

```

y_star = model(x_star.unsqueeze(0)).argmax(dim=1)
Y_prime = 1 - y_star # binary case

X_obs = dataset.tensors[0].to(device)
weights = torch.ones(len(X_obs), device=device) / len(X_obs)

problem = make_cf_problem(model, x_star, Y_prime, X_obs, weights)

algorithm = NSGA2(pop_size=200)
termination = get_termination("n_gen", 150)

# no seed => nondeterministic search
res = minimize(problem, algorithm, termination, verbose=True)

F_mmo, X_mmo = res.F, res.X

```

n_gen	n_eval	n_nds	eps	indicator
1	200	40	-	-
2	400	60	0.0074513698	ideal
3	600	94	0.0513458059	nadir
4	800	150	0.0451504967	nadir
5	1000	200	0.0311961036	nadir
6	1200	200	0.0162890534	nadir
7	1400	200	0.0075316799	ideal
8	1600	200	0.0118130456	nadir
9	1800	200	0.0055258211	nadir
10	2000	200	0.0034698238	f
11	2200	200	0.0065894592	nadir
12	2400	200	0.5000000000	ideal
13	2600	200	0.0108915665	nadir
14	2800	200	0.0034307193	f
15	3000	200	0.0066605826	f
16	3200	200	0.0169677274	nadir
17	3400	200	0.0077844034	nadir
18	3600	200	0.0043747161	f
19	3800	200	0.0046671153	f
20	4000	200	0.0078454757	nadir
21	4200	200	0.0048124385	f
22	4400	200	0.0033250719	f
23	4600	200	0.0052012007	f
24	4800	200	0.0047890662	f
25	5000	200	0.0049239703	f
26	5200	200	0.0028575481	nadir
27	5400	200	0.0071292891	nadir
28	5600	200	0.0052347535	f
29	5800	200	0.0033653114	f

30	6000	200	0.0037342471	f
31	6200	200	0.0049235177	nadir
32	6400	200	0.0047452207	f
33	6600	200	0.0053395533	f
34	6800	200	0.0120562279	nadir
35	7000	200	0.0034236133	f
36	7200	200	0.0039912565	f
37	7400	200	0.0048216101	f
38	7600	200	0.0031445556	f
39	7800	200	0.0126888583	nadir
40	8000	200	0.0039170790	f
41	8200	200	0.0044806463	f
42	8400	200	0.0036271695	f
43	8600	200	0.0032980388	f
44	8800	200	0.0045974664	f
45	9000	200	0.0128973900	nadir
46	9200	200	0.0128586833	nadir
47	9400	200	0.0065780469	f
48	9600	200	0.0041116870	f
49	9800	200	0.0053098939	f
50	10000	200	0.0048270585	f
51	10200	200	0.0035304258	f
52	10400	200	0.0046016040	nadir
53	10600	200	0.0046228766	nadir
54	10800	200	0.0048943296	f
55	11000	200	0.0048947933	f
56	11200	200	0.0036914294	f
57	11400	200	0.0033226543	f
58	11600	200	0.0037516508	f
59	11800	200	0.0040519822	f
60	12000	200	0.0052028189	f
61	12200	200	0.0029925912	f
62	12400	200	0.0032377956	f
63	12600	200	0.0036506532	f
64	12800	200	0.0044730447	f
65	13000	200	0.0042951742	f
66	13200	200	0.0043840924	f
67	13400	200	0.0038819299	f
68	13600	200	0.0037165446	f
69	13800	200	0.0049938553	f
70	14000	200	0.0031738040	f
71	14200	200	0.0047548071	f
72	14400	200	0.0044246227	f
73	14600	200	0.0035528159	f
74	14800	200	0.0048461517	f
75	15000	200	0.0034655119	f
76	15200	200	0.0023737388	f
77	15400	200	0.0051952969	f

78	15600	200	0.0029768615	f
79	15800	200	0.0036218609	f
80	16000	200	0.0027665731	f
81	16200	200	0.0040260739	f
82	16400	200	0.0044229637	f
83	16600	200	0.0036679947	f
84	16800	200	0.0041204997	f
85	17000	200	0.0047870921	f
86	17200	200	0.0031880090	f
87	17400	200	0.0043575809	f
88	17600	200	0.0045959817	f
89	17800	200	0.0045771937	f
90	18000	200	0.0040913895	f
91	18200	200	0.0044173212	f
92	18400	200	0.0049210327	f
93	18600	200	0.0035746081	f
94	18800	200	0.0031885218	f
95	19000	200	0.0047175058	f
96	19200	200	0.0038728401	f
97	19400	200	0.0035890833	f
98	19600	200	0.0061153614	f
99	19800	200	0.0036474702	f
100	20000	200	0.0040699498	f
101	20200	200	0.0038524079	f
102	20400	200	0.0034879370	f
103	20600	200	0.0049641155	f
104	20800	200	0.0041284756	f
105	21000	200	0.0031381285	f
106	21200	200	0.0054959466	f
107	21400	200	0.0038063753	f
108	21600	200	0.0028959520	f
109	21800	200	0.0043060078	f
110	22000	200	0.0026290482	f
111	22200	200	0.0038648687	f
112	22400	200	0.0034890358	f
113	22600	200	0.0040045428	f
114	22800	200	0.0045793384	f
115	23000	200	0.0049557637	f
116	23200	200	0.0024040754	f
117	23400	200	0.0083459633	f
118	23600	200	0.0043890618	f
119	23800	200	0.0036891827	f
120	24000	200	0.0038066902	f
121	24200	200	0.0052014247	f
122	24400	200	0.0063508177	f
123	24600	200	0.0044058182	f
124	24800	200	0.0034870874	f
125	25000	200	0.0044639002	f

126		25200		200		0.0043051905		f
127		25400		200		0.0038890054		f
128		25600		200		0.0045799123		f
129		25800		200		0.0038052106		f
130		26000		200		0.0042351763		f
131		26200		200		0.0055448765		f
132		26400		200		0.0037815549		f
133		26600		200		0.0036317236		f
134		26800		200		0.0042728158		f
135		27000		200		0.0034277282		f
136		27200		200		0.0040730193		f
137		27400		200		0.0040445394		f
138		27600		200		0.0043772804		f
139		27800		200		0.0039593161		f
140		28000		200		0.0041312695		f
141		28200		200		0.0040233655		f
142		28400		200		0.0046377608		f
143		28600		200		0.0055924499		f
144		28800		200		0.0048658302		f
145		29000		200		0.0033648750		f
146		29200		200		0.0038273163		f
147		29400		200		0.0037634887		f
148		29600		200		0.0071680760		nadir
149		29800		200		0.0033554744		f
150		30000		200		0.0034740292		f

```
[115]: # -----
# 4) Post-processing & plots
# -----
# With prob-based validity, use threshold instead of ==0.
# valid_mask = F_mmo[:, 0] < 0.05 # e.g. at least 95% target prob
# valid_F = F_mmo[valid_mask]
# valid_X = X_mmo[valid_mask]

# plt.figure()
# plt.scatter(X[:, 0], X[:, 1], c=Y, label="Original data", s=15, 
#             cmap="coolwarm")
# plt.scatter(X_mmo[:, 0], X_mmo[:, 1], c="red", marker="x", label="All CFs")
# plt.scatter(valid_X[:, 0], valid_X[:, 1], c="black", marker="x", label="Valid CFs")
# plt.scatter(x_star[0].item(), x_star[1].item(),
#             c="green", marker="*", s=200, label="x*")
# plt.legend()
# plt.show()

fig,ax=plot_torch_proba_boundary(model, X)
```


```

ax.scatter(X_mmo[:, 0], X_mmo[:, 1], c="red", marker="x", label="All CFs")
ax.scatter(x_star[0].item(), x_star[1].item(),
           c="green", marker="*", s=200, label="x*")
ax.legend()
plt.show()

```

C:\Users\yuanl\AppData\Local\Temp\ipykernel_26632\3152812757.py:33:
MatplotlibDeprecationWarning:

The `get_cmap` function was deprecated in Matplotlib 3.7 and will be removed in 3.11. Use `matplotlib.colormaps[name]` or `matplotlib.colormaps.get_cmap()` or `pyplot.get_cmap()` instead.



```


[116]: validated_F_mmo=np.array([f for f in F_mmo if f[0]<=0.5])
validated_X_mmo=np.array([X_mmo[i] for i in range(len(F_mmo)) if F_mmo[i][0]<=0.
                           ↪5])
not_validated_F_mmo=np.array([f for f in F_mmo if f[0]>=0.5])
not_validated_X_mmo=np.array([X_mmo[i] for i in range(len(F_mmo)) if ↪
                           ↪F_mmo[i][0]>=0.5])

```

```
valided_F_mmo.shape, valided_X_mmo.shape, not_valided_F_mmo.  
↳shape, not_valided_X_mmo.shape
```

```
[116]: ((144, 4), (144, 2), (56, 4), (56, 2))
```

```
[117]: # visualize the Pareto front in objective space  
fig, ax = plt.subplots(2, 3, figsize=(15, 10))  
  
# Create pairs of objective indices to plot  
obj_pairs = [(1, 3), (2, 3), (1, 2), (0, 1), (0, 2), (0, 3)]  
obj_labels = ['Validity', 'Similarity', 'Sparsity', 'Plausibility']  
  
for i, (obj_x, obj_y) in enumerate(obj_pairs):  
    row = i // 3  
    col = i % 3  
    ax[row, col].scatter(problem.evaluate(X)[:, obj_x],  
                         problem.evaluate(X)[:, obj_y],  
                         marker='*', c='green', s=200, label='Original Instance  
↳x* in Obj Space')  
    ax[row, col].scatter(F_mmo[:, obj_x], F_mmo[:, obj_y], c='blue',  
↳label='Pareto Front')  
    ax[row, col].scatter(valided_F_mmo[:, obj_x], valided_F_mmo[:, obj_y],  
↳c='black', label='Valid Counterfactuals')  
    ax[row, col].set_xlabel(obj_labels[obj_x])  
    ax[row, col].set_ylabel(obj_labels[obj_y])  
    ax[row, col].legend()  
  
plt.tight_layout()  
plt.show()
```



```
[118]: import plotly.graph_objects as go

F_X = problem.evaluate(X)

fig = go.Figure()

fig.add_trace(go.Scatter3d(
    x=F_X[:,1], y=F_X[:,2], z=F_X[:,3],
    mode='markers',
    marker=dict(color='green', size=6, opacity=0.2),
    name='Original Instance x* in Obj Space'
))

fig.add_trace(go.Scatter3d(
    x=not_valided_F_mmo[:,1], y=not_valided_F_mmo[:,2], z=not_valided_F_mmo[:,3],
    mode='markers',
    marker=dict(color='blue', size=3, opacity=0.2, symbol='x'),
    name='Pareto Front which not valid'
))

fig.add_trace(go.Scatter3d(
    x=valided_F_mmo[:,1], y=valided_F_mmo[:,2], z=valided_F_mmo[:,3],
    mode='markers',
    marker=dict(color='black', size=3, opacity=0.2),
    name='Valid Counterfactuals'
))
```

```
    mode='markers',
    marker=dict(color='red', size=5, symbol='cross'),
    name='Valid Counterfactuals in Obj Space'
))

fig.update_layout(
    scene=dict(
        xaxis_title='similarity/AU',
        yaxis_title='sparsity',
        zaxis_title='plausibility'
    ),
    width=900, height=700
)

fig.show()

#      html
fig.write_html("pareto_front_3d.html")
```

[]: